

AI Calling Strategy

Ditto Insurance — Product Strategy Case Study

Scope Stages 2–3 of the sales funnel (lead qualification + slot confirmation)

Platform Bolna AI (pilot) + Exotel (scale)


Pilot 100 inbound leads | 30 days | non-DND | opted-in

Optimising for Deployment speed + brand safety

Target 35%+ slot-book rate at <40% of human advisory cost per slot

One-line thesis. Ditto's only constraint is trust — mis-selling risk, spam perception, and advisory quality. AI should automate the two stages where none of that risk exists (logistics, scheduling), and stay entirely out of the stages where it does (advice, disclosure, underwriting). This document is built around that boundary.

Customer Journey at a Glance



Red boxes = AI scope (scheduling only). Navy boxes = human-handled. Pain callouts show where Ditto loses 85% of leads before the first advisor touchpoint.

SECTION A

Sales Funnel & AI Opportunity

Q1 — Mapping the Funnel

Ditto's funnel spans seven stages from awareness to policy issuance. Conversion estimates below are drawn from three sources: the LeadSquared Insurance Digital Survey (India B2C), Unbounce Finance & Insurance Benchmarks 2024, and the Zopper India InsurTech Report. Where ranges are given, I've used the conservative end for a trust-first brand like Ditto.

Stage	What Happens	Drop-off Driver	Conv. Rate*	AI Fit
1 — Awareness	Visitor arrives via Finshots, SEO, referral	Not intent-driven; high bounce	100% → 15%	None
2 — Lead Qual ★	Fills form; selects callback slot	Form fatigue; no slot availability; trust deficit	15% → 8%	PRIMARY
3 — Confirmation ★	Slot confirmed; reminder sent	No-shows; spam perception; change of mind	8% → 5%	PRIMARY
4 — Advisory Call	30-min IRDAI advisor consultation	Advisor unavailable; not ready to commit	5% → 3.5%	Human only
5 — Follow-Up	WhatsApp summary + policy nudges	Advisor bandwidth; evaluating alternatives	3.5% → 2%	Phase 2
6–7 — Decision & Issuance	Policy chosen, underwriting, issued	Disclosure confusion; medical rejections	2% → 1.2%	Human only

★ Stages 2–3 are the primary AI opportunity. * Sources: LeadSquared Insurance Digital Survey; Unbounce Finance & Insurance Benchmarks 2024; Zopper India InsurTech Report.

Why Stages 2–3, and not others?

The prioritisation logic has three layers:

Zero advice required. These stages are purely logistical — confirming intent and booking a slot. No insurance knowledge changes hands, which means zero mis-selling risk and no IRDAI intermediary exposure. The script can be reviewed and locked by legal once, not call-by-call.

Largest volume loss. Over 85% of leads are lost before Stage 4. Even a modest 10-percentage-point improvement in Stages 2–3 conversion adds ~150 monthly consultations without increasing advisor headcount.

Speed-to-call is a proven lever. LeadSquared data shows that calling within 5 minutes of form submission can lift conversion by 50%+. Human advisors cannot consistently hit that SLA at scale. An AI caller can.

Stage 4 is Ditto's moat. The 30-minute advisory call is irreplaceable — it is the product. Automating it would not just risk mis-selling; it would directly undermine the brand positioning that earns Ditto its 4.9 rating. That boundary is non-negotiable in this proposal.

A note on validation. The funnel analysis above is built on industry benchmarks — not Ditto-specific data. Before scaling beyond the pilot, I'd interview 5–8 Ditto customers on whether an AI scheduling call would feel helpful or intrusive, and A/B test AI call vs. WhatsApp-first for lead confirmation. The pilot should be designed to generate that signal, not just hit a slot-book rate target.

User Segmentation Hypothesis

Not all Ditto leads behave the same way. Before designing a single AI calling flow, it's worth stating the segmentation hypothesis the pilot should test:

Segment	Likely Behaviour	Preferred Channel	Pilot Design Implication	Success Metric
Under 35 (Urban, digital-native)	Comfortable with async. May find an unsolicited call intrusive if they expected WhatsApp. Higher form-abandonment if forced into a call flow.	WhatsApp-first → call only if link not clicked within 30 min	A/B test: Arya call vs. WhatsApp booking link as first touchpoint. Track slot-book rate by segment.	Slot-book rate (did they book at all, and via which channel?)
35–50 (Working professional)	Familiar with advisory calls. More likely to trust a warm voice confirmation. Less likely to respond to async nudges.	AI call within 5 min → WhatsApp confirmation	Primary target for the pilot AI call flow. Highest ROI segment for voice.	Show-up rate (did they actually attend the booked slot?)
50+ (Pre-retirement, family planner)	May be confused by AI caller if not identified clearly. Likely to ask insurance questions Arya cannot answer. Higher escalation risk.	AI call with immediate human escalation path	Flag for faster warm-transfer SLA. Review call transcripts weekly for mis-escalation patterns.	Escalation satisfaction (post-transfer advisor rating; did the handoff feel seamless?)

These are hypotheses, not conclusions. The pilot's call outcome data — segmented by age and city tier — should validate or overturn them within 30 days.

Platform Research & Recommendation

Q2 — Outreach & Comparison

I contacted all three platforms directly rather than relying solely on published pricing pages, since voice AI rates in India vary significantly by volume tier and often go unpublished. Bolna AI responded via live chat within 4 hours (transcript saved). Plivo pricing was verified on their India voice rates page and confirmed in a follow-up sales chat. Exotel was reached via their enterprise enquiry form; an account executive confirmed the credit-pack entry point but declined to share a per-minute rate without a volume commitment — reflected as an estimate below.

Criteria	Plivo	Bolna AI	Exotel
Pricing (India)	■0.74/min + ASR. Eff. ■8–12/min all-in	■7/min standard; ■3/min at 500+ min/day. ~\$0.06 all-in incl. LLM	Credit-based; ■9,999 entry. Est. ■8–15/min (no pub. rate)
AI Voice Capability	CPaaS + APIs only. No native agent — custom AI layer needed.	Full voice AI (STT + LLM + TTS). 20+ models. <300ms latency.	Voicebot via Cogno AI (acq.). Bidirectional streaming. Not core product.
Indian Language Support	Hindi, English. Limited Hinglish.	10+ languages: Hinglish, Tamil, Telugu, Bengali, Marathi.	5+ languages. Strong Hindi. Pan-India telco coverage.
Integration & Compliance	REST API + webhooks. Manual DND filtering required.	Native CRM + built-in NDNC check. Auto calling-hour enforcement. DPDPA-ready.	Pre-built Salesforce, LeadSquared. Full TRAI compliance. India data residency. Best for BFSI.
Setup Complexity	Moderate-High. Build AI layer separately. 2–4 weeks.	Low. No-code builder. Deploy in hours to days.	Moderate-High. Enterprise. Longer sales cycle.
Key Limitation	Not a full voice agent — significant engineering overhead.	Newer platform (~2024). Smaller enterprise track record.	AI voicebot is add-on, not core. Pricing opaque without commitment.

Recommendation

Pilot: Bolna AI — the only platform that combines a full native voice agent, Hinglish support, transparent pricing, and a no-code builder. For a 30-day pilot with 100 leads, the priority is deployment speed and brand safety, not enterprise-grade audit trails. Bolna's <300ms latency also matters beyond specs: pauses over 700ms consistently trigger 'hello?' responses that signal 'robot' to the caller and break script flow. I'd treat p95 latency as a trust KPI in the pilot, not just a technical one.

Scale (5,000+ calls/month): Exotel — BFSI compliance depth and full India data residency justify the longer onboarding when volume is there. Not for a 100-lead pilot.

Not now: Plivo — excellent raw telephony at ■0.74/min, but it's infrastructure, not a turnkey agent. Would add 3–5 weeks of engineering work that the pilot doesn't justify.

SECTION C

AI Caller Design — Arya

Q3 — Persona

The persona is not a branding exercise — it is a constraint document. Every attribute below defines something Arya will not do as much as what she will.

Name	Arya
Role	Ditto's scheduling and confirmation assistant. Not a sales agent. Not an advisor.
Voice	Clear, warm, unhurried. Hinglish by default; full Hindi or English on switch. Never robotic, never pushy.
Personality	Helpful and matter-of-fact. Transparent about being AI when asked. Never attempts insurance guidance.
Brand Alignment	Embodies Ditto's no-spam promise. Never asks twice if someone says no. Never upsells. Logistics only.
Tone Restrictions	No urgency language ('limited slots', 'offer ends tonight'). No jargon. No pressure. Customer always in control.

Q3 — Call Script: Stage 2 Slot Booking (Priority 1)

Trigger: Lead form submitted. AI calls within 5 minutes. Calling hours: 10 AM–7 PM. Target speaking time: under 55 seconds.

Branch	Arya's Script
Greeting	"Hello, am I speaking with [Name]? This is Arya from Ditto Insurance — you just filled our enquiry form. I'm Ditto's scheduling assistant, this is a 30-second call, not a sales pitch. Is now okay?"
Customer agrees	"Great! Our IRDAI-certified advisors offer free, no-pressure consultations. Most customers prefer a morning slot — would 11 AM tomorrow work, or is afternoon better?"
Hesitation	"Of course, no pressure! One thing though — our advisors never sell on the first call. It's genuinely a 30-minute conversation. Shall I send a WhatsApp first so you can read what the call covers before deciding?"
Slot confirmed	"Perfect! I've booked [Day] at [Time]. You'll get a WhatsApp confirmation shortly. Zero spam, zero pressure from Ditto's side."
Soft intel (optional)	"Anything specific you'd like the advisor prepared for — like a pre-existing condition or a budget range? Great, I've noted that. Thank you, [Name]!"
Not a good time	"No problem! I'll send a WhatsApp link to pick a time that suits you. Take your time — there's no rush. Goodbye!"
Are you a robot?	"Yes, I'm Arya, Ditto's AI assistant — I only handle scheduling. All insurance advice comes from real, IRDAI-certified advisors. Want me to book a slot?"

Design note: 'Most customers prefer a morning slot' is deliberate social proof anchoring. It reduces cognitive load vs. an open calendar prompt without being manipulative — the customer still chooses the time.

Q4 — Edge Cases

The 23 edge cases below are organised across four domains. The handling logic is consistent with one principle: when in doubt, exit gracefully and hand off to WhatsApp or a human. An AI that keeps talking when it shouldn't is a brand liability.

Customer Behaviour

#	Edge Case	Handling Flow	Why It Matters
1	Busy right now	Apologise, send WhatsApp booking link, end call. No retry for 4+ hours. Max 1 retry per lead.	Ditto's no-spam promise is non-negotiable.
2	No answer	No voicemail. Mark no-answer. Retry twice: +2 hrs, next morning. After 3 attempts — WhatsApp-only.	Repeated calls = spam perception.
3	Voicemail / IVR	AMD detects in 3s. No message left. Logs attempt; switches to WhatsApp.	AI insurance voicemail feels spammy.
4	Already booked	Check CRM in real-time. If slot found — confirm gracefully. If no record — log discrepancy, route to ops.	Duplicate calls break trust.
5	Asks insurance question	Hard block on policy guidance. Arya: 'Great question for our advisor — shall I book a call?'	AI insurance advice = mis-selling risk + IRDAI liability.
6	Angry / frustrated	De-escalate. 'I'm sorry for the inconvenience. Shall I remove your number?' Log negative sentiment. Human notified.	Respect is core to Ditto's brand.
7	'Remove me from list'	Confirm opt-out, thank, end. Flag DNC in CRM in real-time.	TRAI compliance + brand integrity.
8	Wants to change slot mid-call	Offer 2 alternate times. If none — send self-scheduling WhatsApp link.	Flexibility increases show-up rate.

Technical Failures

#	Edge Case	Handling Flow	Why It Matters
1	High latency (>2s)	Play filler ('Just a moment...'). If persistent — end gracefully; log; retry later.	Long pauses cause hang-ups; damage brand.
2	ASR misrecognition	If confidence below threshold — confirm: 'Just to confirm, you said [X]?' After 2 fails — WhatsApp fallback.	Wrong slot bookings waste advisor time.
3	Call drops mid-conversation	CRM logs drop + last utterance. Auto WhatsApp: 'Our call got disconnected. Here's a link to reschedule.'	Drop without follow-up = lost lead.
4	CRM API unavailable	Proceed with soft caveat; log for manual reconciliation. Don't block calls on CRM downtime.	Single point of failure unacceptable.
5	Platform outage (Bolna / Exotel)	Fallback to WhatsApp-first queue. Ops alerted via Slack. High-priority leads handled manually.	Resilience is non-negotiable.

Language & Conversation

#	Edge Case	Handling Flow	Why It Matters
1	Switches to Hindi mid-call	Bolna mirrors language blend. Full Hindi switch — continue in Hindi. Preference logged for advisor.	Non-adaptive AI feels foreign.
2	Unsupported regional language	Arya: 'Let me send a WhatsApp instead — our team can help in your preferred language.' Route to regional queue.	Honesty beats false fluency.

#	Edge Case	Handling Flow	Why It Matters
3	Goes off-script (long personal story)	Listen 15–20s, then: 'That sounds difficult. Our advisors handle exactly this — shall I set up a call?'	Empathy without derailing = Ditto brand.
4	Long silence (>8s)	Arya: 'Are you still there?' After another 8s — 'No problem, I'll send a message. Goodbye!'	Hanging silently wastes resources.

Escalation Triggers

#	Edge Case	Handling Flow	Why It Matters
1	Explicitly asks for a human	Warm transfer to available advisor. If none, schedule callback within 30 min.	Never lock a customer in an AI loop.
2	Expresses distress or grief	Pause script. 'I'd like to connect you with an advisor directly.' Flag URGENT in CRM.	Human empathy irreplaceable in vulnerable moments.
3	Mentions prior insurer decline	Arya: 'Our advisors specialise in this. Let me get you connected.' Flag 'Prior decline — specialist needed.'	Wrong guidance = serious financial consequences.
4	DND-registered number called	NDNC pre-call check via TRAI API. If on DND — call blocked; WhatsApp sent (transactional consent only).	TRAI violations carry significant penalties.
5	Internal advisor number in queue	CRM flags number as internal. System suppresses call. Log as INTERNAL_NUMBER. Alert ops.	AI calling an advisor is a data quality failure.
6	Customer gives different callback number mid-call	Capture via ASR. Flag for human ops to verify before updating CRM. Don't auto-overwrite. Send WhatsApp to original number.	ASR transcription of phone numbers is error-prone. Auto-update risks corrupting the lead record.

SECTION D

Execution Plan, Data & KPIs

Q5 — 6-Week Execution Plan

Week	Phase	Key Steps	Owners / Dependencies
Wk 1	Discovery & Alignment	CEO sign-off on pilot scope. Define success criteria. Map current lead flow in CRM. Audit existing infra. Define legal scope: TRAI, IRDAI, DPDPA.	CEO, Product, Sales Ops, Legal
Wk 1–2	Vendor Outreach	Formally engage Bolna AI + Exotel for pilot pricing. Request sandbox access. Negotiate: 30-day pilot, volume cap, data residency SLA. Sign NDAs + DPAs.	Product, Engineering, Legal
Wk 2	Script Development	Draft call scripts (Slot Booking + Reminder). Internal review with advisors. Legal review: no language construable as insurance advice. A/B test Hinglish vs. English voice.	Product, Advisors, Legal, Brand
Wk 2–3	Technical Integration	Set up Bolna workspace + agent config. Configure Exotel: +91 numbers, DND filtering, routing webhooks. Build CRM integration. Implement fallback logic.	Engineering (1–2 devs), DevOps
Wk 3	Internal Testing	Test all 23 edge cases manually. QA: ASR accuracy on Hindi/Hinglish. Stress-test 20 simultaneous calls. Verify DND filter. Go/No-Go decision.	Engineering, QA, 2 Advisors
Wk 4	Soft Pilot — 100 Leads	100 fresh inbound leads (non-DND, opted-in). Slot Booking flow only. Calling hours: 10 AM–7 PM. Monitor daily. Collect advisor feedback. Survey 5–10 customers.	Product Lead daily, Advisors
Wk 5–6	Iteration & Expansion	Review KPIs. V2 script update. Add Reminder flow (Priority 2). Expand: 100 → 500 → full pipeline. Contingency: no live calls before legal sign-off.	Product, Engineering, Sales Ops

Time-of-Day Calling Strategy

Windows below are based on pickup-rate patterns for working professionals in India, and are deliberately conservative relative to TRAI's permitted 9 AM–9 PM window. The extra margin is intentional — the risk of one frustrated 9 PM insurance call outweighs the marginal conversion gain.

Window	Time Slot	Rationale	Priority
Morning (best)	10 AM – 12 PM	Post-commute, pre-lunch. Highest pickup intent.	Priority 1
Afternoon (good)	2 PM – 5 PM	Post-lunch settled, pre-evening commute.	Priority 2
Saturday	11 AM – 2 PM only	Weekend receptivity window. Stop by 2 PM.	Weekend only
Evening (avoid)	7 PM – 9 PM	Family time. Higher frustration even if TRAI-compliant.	Avoid
Morning rush (avoid)	8 AM – 10 AM	Commute window. Low pickup, high irritation.	Avoid
Lunch (avoid)	12 PM – 2 PM	Low availability. Distracted context.	Avoid

Optimise calling windows weekly using pilot pickup-rate data. If morning pickup rate exceeds afternoon by 10%+ for 7 consecutive days, deprioritise afternoon entirely.

Q6 — Data Formats

INPUT — JSON webhook from CRM on lead creation

Field	Type	Description	Example
lead_id	String (UUID)	Unique identifier from CRM	lead_20260417_a3f9
full_name	String	Customer name for greeting personalisation	Priya Sharma
phone_number	String (+91)	Mobile. E.164. Pre-checked vs. NDNC.	+919876543210
insurance_type	Enum	From form field. Personalises opening.	health term both
preferred_language	Enum	Defaults to hinglish if missing.	hi en hinglish
city	String	Time-zone and cultural context.	Mumbai
form_submit_time	ISO 8601	Triggers call scheduling.	2026-04-17T11:32:00+05:30
dnd_status	Boolean	Pre-checked via NDNC API. True = no call.	false
advisor_notes	String (opt.)	Pre-existing context from form comments.	Father has Type 2 diabetes
available_slots	Array of ISO timestamps	Next 5 open advisor slots for Arya to offer.	['2026-04-18T11:00+05:30',...]
session_engagement_score	Integer (0–100)	Time-on-site + pages visited. Powers call priority.	78 — call within 3 min if >70

OUTPUT — logged after each call

Field	Type	Description	Example
call_id	String (UUID)	Unique call ID from Bolna/Exotel.	call_0x9a2f...
lead_id	String	Links back to originating lead record.	lead_20260417_a3f9
call_timestamp	ISO 8601	When the call was placed.	2026-04-17T11:37:00+05:30
call_duration_secs	Integer	Total call duration in seconds.	47
call_outcome	Enum	Primary result of the call.	SLOT_BOOKED CALLBACK_REQUESTED DNC_OPT_OUT NO_ANSWER VOICEMAIL ESCALATED_TO_HUMAN CALL_FAILED
slot_booked_time	ISO 8601 (nullable)	Populated if SLOT_BOOKED.	2026-04-18T11:00+05:30
customer_sentiment	Enum	AI sentiment analysis of call.	POSITIVE NEUTRAL NEGATIVE DISTRESSED
escalation_flag	Boolean	True if human takeover triggered mid-call.	false
retry_count	Integer	Which attempt number this was (1, 2, or 3).	1
transcript_summary	String	3-sentence AI summary for advisor context.	Customer confirmed slot 11 AM. Asked about parent cover. No existing policy.
latency_ms	Integer	Bolna response latency. Flag if >700ms.	280

Q7 — KPIs

Five KPIs, each tied to a specific decision it should trigger. Targets are directional for pilot sizing — calibrate actuals from Wk 4 data before committing to scale-stage benchmarks.

KPI	Pilot Target	How Measured	Why It Matters & How Target Was Set
1 — Slot Book Rate	≥35%	Slots booked / Leads reached	Primary success metric. Target derived from LeadSquared insurance benchmarks (20–25% for cold outbound) with a +10pp uplift assumed for warm inbound leads who just submitted a form. Below 35% = script needs revision before scale.
2 — Connect Rate	≥45%	Calls answered / Calls placed	Below 30% signals caller-ID reputation or wrong calling hours — both fixable. 45% is the midpoint of industry benchmarks for B2C voice in India (35–60% depending on channel and list quality). Tracked daily in the pilot.
3 — Show-Up Rate	≥65%	Consultations completed / Slots booked by AI	Validates that AI-confirmed slots are as strong as human-confirmed. Ditto's current human show-up rate (estimated 70–75% per advisor feedback) sets the baseline. AI at ≥65% = acceptable. Below 55% = trust gap in AI confirmation flow.
4 — Complaint & Escalation Rate	<1%	DNC opt-outs + negative sentiment / Total calls	Guardrail KPI. 2% spike = immediate pilot pause. 1% threshold based on TRAI complaint norms and Ditto's reputational standard. One viral complaint about Ditto's 'spam AI calls' undoes months of trust-building.
5 — Cost per Slot vs. Human Baseline	<40% of human cost	AI cost per slot / Human advisor cost per slot	Human advisor opportunity cost per slot booked: estimated ₹150–200 (30 min of advisor time at median compensation, even for purely logistical calls). AI target: ₹40–60 at ₹7/min Bolna rate with ~6-min average call. Below 40% = business case for scale.

On the numbers: KPI targets above are anchored to published industry benchmarks and internally consistent logic, not guesses. They should be validated against Ditto's actual CRM data in Week 1 of the execution plan — if Ditto's current human show-up rate is 80%, the AI show-up target should be recalibrated upward accordingly.

SECTION E

Dashboard & Monitoring

Q8 — Sales Manager Dashboard

The dashboard is designed to answer one question in under 60 seconds each morning: *Did the AI perform, are leads stuck, and are there any brand risk signals?* Every element drives a decision — not just a data point.

Zone	What It Shows	Purpose & Action It Drives
Zone 1 — Daily Health Strip (full width, top; refreshes every 15 min)	5 tiles: Leads Called Connect % Slots Booked Escalations Complaints	Daily health at a glance. Complaints tile turns red if >2 today. Each tile shows today's value + 7-day sparkline. Click opens filtered drill-down list.
Zone 2 — 30-Day Funnel Chart (left column)	Rolling 30-day funnel: Leads Called → Connected → Slots Booked → Consulted → Policy. Secondary line shows human-confirmed funnel for comparison.	Identifies where volume leaks. Each stage shows absolute count + % drop from stage above. Clicking a stage filters the live feed to that outcome only.
Zone 3 — Live Call Feed (centre column)	Last 20 calls: Lead Name Insurance Type Outcome Badge Sentiment Icon Timestamp. Click to expand transcript + recording link.	Real-time visibility. NEGATIVE / DISTRESSED flags show one-click 'Assign Advisor' button that creates a priority CRM task. DISTRESSED rows highlighted in red.
Zone 4 — Cost & Efficiency Panel (right column)	AI spend today (■) Cost/slot vs. ■150 human baseline Advisor hours saved Bolna p95 latency. 7-day bar chart: AI vs. human show-up rate. A/B toggle: Hinglish vs. English.	Makes the business case visible daily. Shows ROI in real time. Latency alert if p95 >700ms.
Zone 5 — Alerts & Action Queue (full width, bottom)	Rule-triggered alert cards: severity icon description recommended action one-click action button. Sorted by severity.	Decision layer. Every alert prescribes an action, not just a data point. Critical alerts (DISTRESSED, complaint rate >1%) push Slack notification to Product Lead within 60 seconds.

Sample Alert Triggers

Trigger	Severity	Recommended Action
Connect rate <35%	Warning	Auto-flag: check calling hours or caller-ID reputation. Switch to 10 AM–12 PM block only.
Complaint rate hits 1%	Critical	Pilot pause triggered automatically. Slack alert to Product Lead within 60 seconds.
Latency p95 >700ms	Warning	Check Bolna server region. Consider fallback to WhatsApp for that hour's lead queue.
DISTRESSED sentiment flagged	Critical	Assign senior advisor immediately. Log for post-pilot UX review.
AI show-up rate > human for 7 consecutive days	Positive	Surface 'Build expansion case' prompt. Present KPI data to CEO for Phase 2 sign-off.

Access: Web dashboard (Sales Manager + senior advisors) | 8 AM daily email digest (5 metrics + top 3 alerts) | Slack real-time alerts for DISTRESSED and ESCALATION calls | Mobile: Zone 1 + Zone 5 only.

Trade-offs, What I Didn't Do & What I'd Change

Key Trade-offs

Advisory calls excluded from AI scope. The risk of reputational damage from AI mis-advice is not worth the trade-off for a brand whose entire moat is trust. Automating Stage 4 would be a faster answer — and the wrong one.

Bolna (speed-to-pilot) over Exotel (maximum compliance). A 30-day pilot does not justify Exotel's enterprise onboarding timeline. Revisit at 5,000+ calls/month when BFSI audit trails become mandatory.

Conservative calling hours: 10 AM–7 PM vs. TRAI-allowed 9 AM–9 PM. The extra two hours carry higher frustration risk. Better to leave marginal revenue on the table than receive one viral complaint about a 9 PM insurance call.

No AI voicemail. Insurance voicemail is perceived as spam. The downside to trust outweighs the marginal callback lift — this was an easy call.

What I Didn't Do (and Why)

Post-consultation follow-up (Stage 5) excluded from pilot. Too many variables for V1. Add in Phase 2 once the slot-booking flow is validated.

No custom AI stack on Plivo. Ditto likely lacks the engineering bandwidth when Bolna deploys in days. Complexity is a risk in a 30-day pilot.

No full-scale cost modelling. Requires actual pilot data on call durations and conversion rates. The ■40–60 per slot estimate is a directional hypothesis, not a projection.

No claims support or policy renewal AI. Different product area requiring separate compliance analysis. Out of scope.

If I Had More Time

Customer validation before solution design. I'd interview 5–8 Ditto customers on what kind of AI contact they'd find acceptable vs. intrusive, and which channel (call vs. WhatsApp) they trust more for insurance scheduling. My hypothesis is under-35 users prefer async WhatsApp; over-35 users respond better to a brief, warm voice confirmation. The pilot should be designed to test this — not assume it.

A/B test call trigger timing: T+5 min vs. T+20 min post-form. Hypothesis: T+5 wins on slot-book rate. T+20 wins on show-up rate — a committed booker reflects before confirming. Younger users likely prefer WhatsApp-first. Segment-dependent; the pilot data should drive this.

Deeper IRDAI regulatory mapping. IRDAI's 2024 guidelines on digital insurance intermediaries require AI systems to identify themselves on every call, not just when asked. Any call touching policy comparison triggers the insurance solicitation definition — the pilot script avoids this, but it must be formally reviewed by a qualified IRDAI compliance officer before deployment. I've flagged this in the execution plan (Wk 1 legal scope) but haven't mapped it in detail. Circular references should be confirmed with legal before go-live.

Negotiate pilot-specific pricing with Bolna and Exotel. Explore a revenue-share model linked to slot-book rate outcomes. Aligns vendor incentives with Ditto's success metric.

Sources

• joinditto.in (product, pricing, reviews)

- zerodha.com/z-connect (Ditto background)
- Plivo India voice pricing — plivo.com/voice/pricing/in (verified via sales chat)
- Bolna AI — bolna.ai; YourStory Nov 2025 interview (live chat transcript retained)
- Exotel AI voicebot documentation (enterprise enquiry email on file)
- LeadSquared Insurance Digital Survey
- Zopper India InsurTech Report
- Unbounce Finance & Insurance Conversion Benchmarks 2024
- Edesy.in India telephony comparison
- Tabbly.io AI Voice Agents India guide
- IRDAI digital insurance guidelines 2024 (exact circular reference to be confirmed with legal counsel before go-live)